

Spatial-MCP + A2A Integration

Proof-of-Concept Project Plan for GCP Deployment

Version 1.2 | December 2025 | Includes Patient-One Reference Architecture

1. Executive Summary

This document outlines the project plan for deploying your existing spatial-mcp bioinformatics MCP servers within a major US research hospital's GCP environment, integrated with Google's Agent2Agent (A2A) framework. The patient-one reference implementation demonstrates a complete multi-modal cancer analysis workflow that serves as the validation baseline for the hospital deployment.

Key Objectives

- Migrate and adapt existing spatial-mcp servers to client's GCP project
- Implement HIPAA-compliant infrastructure with BAA coverage
- Integrate MCP tool servers with A2A agent orchestration
- Deploy on Vertex AI Agent Engine for production scalability
- Replicate patient-one workflow (PAT001-OVC-2025) as validation baseline
- Deliver audience-specific outputs: developer, care team, and patient

2. Patient-One Reference Implementation

The patient-one directory demonstrates a complete end-to-end multi-modal cancer analysis for a high-grade serous ovarian carcinoma case (Patient ID: PAT001-OVC-2025). This validated workflow serves as both proof-of-concept and acceptance criteria for hospital deployment.

2.1 Reference Patient Case

Attribute	Value
Patient ID	PAT001-OVC-2025 (Sarah Elizabeth Anderson)
Diagnosis	High-grade serous ovarian carcinoma, Stage IV
Key Mutations	BRCA1 germline, TP53 R175H, PIK3CA E545K, PTEN LOH
Resistance Status	Platinum resistant (CA-125: 1456 → 22 → 389 → 289 U/mL)
TCGA Subtype	C1 (Immunoreactive) / C2 (Differentiated)
TME Classification	Immunologically COLD (immune-excluded phenotype)

2.2 Five-Test Integration Suite

The workflow executes five integrated tests, each building upon previous results:

Test	Analysis Type	MCP Servers	Key Outputs
Test 1	Clinical Data & Genomic Analysis	mcp-mockepic, mcp-tcga, mcp-fgbio	Patient demographics, mutations, TCGA classification
Test 2	Multi-Omics Resistance Analysis	mcp-multiomics	Stouffer's Z-scores, pathway activation, resistance mechanisms
Test 3	Spatial Transcriptomics	mcp-spatialtools	900 spots, 6 regions, immune/resistance marker distribution
Test 4	Histology & Imaging	mcp-openimagedata, mcp-deepcell	H&E composition, CD8+ counts, Ki67 index, cell phenotypes
Test 5	Integrated Analysis & Recommendations	All servers (synthesis)	Clinical recommendations, therapeutic targets, prognosis

2.3 Audience-Specific Output Structure

The patient-one-outputs directory organizes deliverables by audience:

for-developer/ (Technical Documentation)

File	Purpose
MCP_Report_PAT001.pdf	Complete technical report documenting all 5 tests, MCP servers used, and data flow
MCP_Servers_Reference_Guide.pdf	Comprehensive documentation of all 10 MCP servers: functions, parameters, use cases
Full_Test_Prompt.pdf	Complete prompts used for each test - enables reproducibility and validation

for-care-team/ (Clinical Outputs)

File	Purpose
MCP_Report_PAT001.pdf	Clinical summary with therapeutic recommendations and prognosis
spatial_transcriptomics_analysis.png	Spatial distribution, gene expression heatmap, proliferation/resistance/immune markers
multomics_resistance_analysis.png	Stouffer's Z-scores, log2 fold changes across RNA/Protein/Phospho, therapeutic targets
histology_imaging_analysis.png	H&E composition (63.6% tumor), CD8+ distribution, Ki67 50%, multiplex IF phenotypes

for-patient/ (Patient-Facing Materials)

File	Purpose
patient_summary.html	Plain-language explanation of diagnosis, treatment plan, and what to expect
medication_guide.html	Medication information, dosing schedules, side effects, and support resources
patient_infographic.png	Visual summary: diagnosis, key findings, treatment plan, timeline, support team

2.4 MCP Server Ecosystem

Ten MCP servers provide the complete bioinformatics toolkit:

Server	Category	Key Functions
mcp-pubmed	Literature	Search articles, get metadata, full-text retrieval from PMC
mcp-fgbio	Genomics	Reference genomes, FASTQ validation, UMI extraction, gene annotations
mcp-huggingface	AI/ML	Genomic foundation models, embeddings, variant effect prediction
mcp-seqera	Pipelines	Nextflow workflow execution, nf-core pipelines, job monitoring
mcp-mockepic	Clinical	EHR access, patient demographics, labs, medications, history
mcp-tcga	Oncology	TCGA cohort comparison, expression data, survival analysis, subtypes
mcp-deepcell	Imaging	Cell segmentation, multiplex IF analysis, phenotype classification
mcp-openimagedata	Imaging	H&E image management, histology storage, annotation
mcp-spatialtools	Spatial	Spatial transcriptomics, spot analysis, region clustering, TME classification
mcp-multiomics	Integration	RNA/Protein/Phospho integration, Stouffer's meta-analysis, HALLA, pathway analysis

2.5 Validation Acceptance Criteria

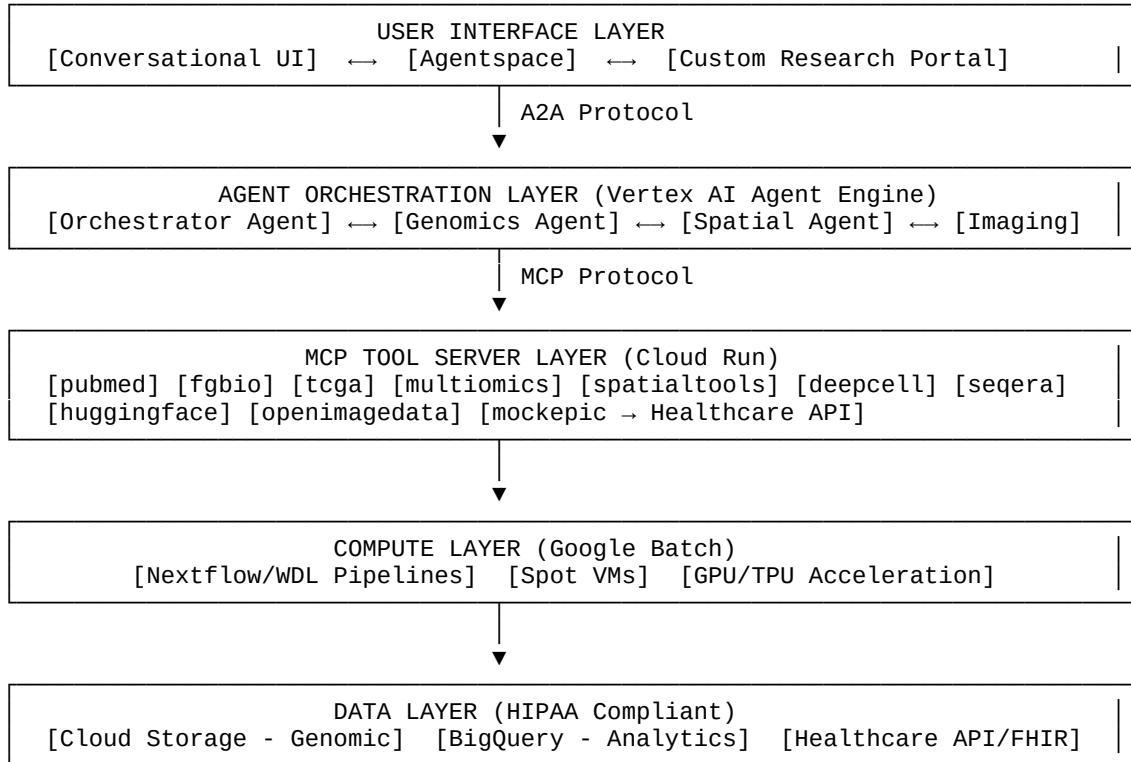
The hospital deployment must replicate these patient-one metrics:

1. **Functional Equivalence:** All 10 MCP servers produce identical outputs given identical inputs
2. **Five-Test Completion:** Orchestrator agent successfully coordinates Tests 1-5 sequence
3. **Audience Outputs:** Generate developer, care-team, and patient deliverables per template
4. **Performance Baseline:** End-to-end workflow within 10% of reference execution time
5. **Statistical Accuracy:** Stouffer's Z-scores, pathway results match within 0.1% tolerance

3. GCP Service Architecture

3.1 Layered Architecture Diagram

The following diagram illustrates the integration of MCP servers with A2A agent orchestration on GCP:



3.2 Core Services Matrix

Component	GCP Service	Purpose
Agent Runtime	Vertex AI Agent Engine	Managed runtime for A2A agents with built-in testing and reliability
MCP Server Hosting	Cloud Run	Serverless container hosting for 10 MCP servers
Foundation Models	Vertex AI (Gemini)	Gemini 2.5 Pro for reasoning and tool use within agents
Pipeline Compute	Google Batch	Nextflow/WDL execution with Spot VMs and GPU acceleration
Data Storage	Cloud Storage / BigQuery	Genomic data lake (GCS) and analytical warehouse (BQ)
Healthcare Data	Healthcare API	FHIR-compliant clinical data (replaces mcp-mockepic in production)

3.3 Google Batch for Bioinformatics Pipelines

Note: Google Batch replaces the deprecated Life Sciences API. Key benefits:

- **Native Nextflow Support:** Direct integration with mcp-seqera workflows
- **GPU/TPU Access:** Accelerated compute for mcp-deepcell segmentation
- **Spot VM Optimization:** Cost-effective for long-running genomic jobs

4. Implementation Phases

Phase 1: Foundation & Compliance (Weeks 1-3)

1. **GCP Project Setup:** Create dedicated project with folder structure
2. **BAA Verification:** Confirm BAA covers Vertex AI, Cloud Run, Google Batch
3. **Network Architecture:** VPC, Private Google Access, VPC Service Controls
4. **IAM & Encryption:** Service accounts, CMEK configuration

Phase 2: MCP Server Migration (Weeks 4-6)

1. **Containerize All 10 Servers:** Docker images with health checks
2. **Cloud Run Deployment:** Deploy each MCP server with appropriate limits
3. **Google Batch Configuration:** Configure for mcp-seqera pipeline execution
4. **Healthcare API Integration:** Replace mcp-mockepic with real FHIR endpoints
5. **Patient-One Validation:** Execute full 5-test suite against reference outputs

Phase 3: A2A Agent Development (Weeks 7-10)

- **ADK Setup:** Install Google Agent Development Kit
- **Create Agent Cards:** A2A AgentCard for each specialized agent
- **Wire MCP to ADK:** Connect 10 MCP servers as agent tools
- **Build Orchestrator:** Primary agent that coordinates 5-test workflow

Phase 4: Agent Engine Deployment (Weeks 11-12)

- **Deploy to Agent Engine:** Push ADK agents to Vertex AI
- **A2A Discovery:** Register /.well-known/agent.json endpoints
- **End-to-End Validation:** Complete patient-one workflow via conversational UI

Phase 5: Validation & Handoff (Weeks 13-14)

- **Output Validation:** Generate all 3 audience deliverables per patient-one template
- **Performance Testing:** Benchmark against patient-one baseline
- **Documentation & Training:** Runbooks, architecture diagrams, knowledge transfer

5. Project Timeline

Phase	Deliverables	Duration	Validation Gate
Phase 1	GCP project, compliance, network	3 weeks	BAA confirmed
Phase 2	10 MCP servers on Cloud Run	3 weeks	Tests 1-4 pass
Phase 3	A2A agents with MCP tools	4 weeks	Test 5 synthesis pass
Phase 4	Agent Engine deployment	2 weeks	Conversational UI works
Phase 5	Validation, docs, training	2 weeks	All 3 outputs generated
TOTAL	Production-ready PoC	14 weeks	Patient-one parity

6. Immediate Next Steps

- Client Kickoff Meeting:** Present plan, walk through patient-one outputs
- BAA Verification:** Confirm coverage for Vertex AI, Cloud Run, Google Batch
- Access Provisioning:** Request IAM roles in client's GCP organization
- Healthcare API Planning:** Coordinate FHIR endpoint access to replace mcp-mockepic
- ADK Familiarization:** Complete ADK tutorials, review A2A spec v0.3

Appendix A: Key Resources

Your Reference Implementation

- spatial-mcp Repository: <https://github.com/lynlangit/spatial-mcp>
- Patient-One Outputs: /architecture/patient-one/patient-one-outputs/
- MCP Servers Reference Guide: for-developer/MCP_Servers_Reference_Guide.pdf

A2A Protocol Resources

- A2A Protocol: <https://a2a-protocol.org/latest/>
- A2A GitHub: <https://github.com/a2aproject/A2A>
- A2A Python SDK: <https://github.com/a2aproject/a2a-python>

GCP Healthcare Resources

- HIPAA Compliance: <https://cloud.google.com/security/compliance/hipaa>
- Google Batch: <https://cloud.google.com/batch/docs>
- Healthcare API: <https://cloud.google.com/healthcare-api/docs>

— *End of Document* —